

Yi's Framework: Structured Loss Analysis Under Time and Order Constraints

Duo Yi*

yiduo2008@gmail.com

Abstract

We present *Yi's Framework*, a operational framework for decision systems that face two ubiquitous structural constraints: *finite-time latency* and *order-sensitivity* (noncommutativity) under reflexive dynamics. We model realized loss as a Bregman regret to an ideal benchmark and prove a *calibrated, additive lower bound*

$$L \geq L_{\text{ideal}} + g_1(\lambda) + g_2(\varepsilon_\star) + g_{12}(\lambda, \varepsilon_\star) - \Delta_{\text{ncx}},$$

Definition (Interaction term g_{12}). Let $a \in \arg \min_{r \in \mathcal{C}_\varepsilon} D_\Phi(p^* \| r)$ and $b \in \arg \min_{r \in \mathcal{C}_\lambda} D_\Phi(a \| r)$ be any measurable selections (selection-free convention). We define

$$g_{12}(\lambda, \varepsilon) := \inf_{q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda} D_\Phi(b \| q).$$

By definition $g_{12} \geq 0$; throughout we assume $\mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda \neq \emptyset$ (feasible intersection).

where g_1 quantifies the latency (time-window) penalty, g_2 quantifies the noncommutative (order) penalty, and $g_{12} \geq 0$ captures their empirical interaction. Here $\Delta_{\text{ncx}} \geq 0$ is a nonconvexity/approximation penalty that *vanishes* under convex or Φ -geodesic-convex conditions; in practice we report the safe nonnegative component $\max\{0, g_1 + g_2 + g_{12} - \Delta_{\text{ncx}}\}$. The bound follows from a two-stage Bregman projection with a generalized Pythagorean inequality. We prove monotonicity, calibration properties, conditions for $g_{12} = 0$ (orthogonality), and give estimation procedures and examples. The result turns diverse domain “impossibilities” into a single geometric statement with actionable diagnostics.

Feasible intersection (Assumption I). We assume $\mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda \neq \emptyset$; otherwise g_{12} is undefined and the bound degenerates (use convexification or redesign).

1 Introduction

Decision systems in products, platforms, science, finance, and policy must act within *finite windows*, wait for *delayed, partial verification*, and operate in *reflexive* environments where actions change the future evidence. Empirically, three tensions recur: (i) one cannot wait forever to verify, (ii) the order of predict/verify/act affects outcomes, (iii) the two tensions can amplify each other.

Goal. We formalize these constraints and prove an *inevitable* additive lower bound on achievable loss, decomposed into a time-like component, an order-like component, and a nonnegative interaction. The proof employs proper scoring rules, Bregman divergences, and two-stage convex projections—yielding a precise and portable calculus.

*Primary contact.

Contributions.

- A unified loss decomposition under time and order constraints.
- A constructive proof via Bregman geometry (two projections + Pythagorean inequality).
- Monotonicity and orthogonality theorems; when and why interaction vanishes.
- Practical diagnostics for estimating g_1, g_2, g_{12} with uncertainty.

2 Setup: Signals, Loss, and Ideals

Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space. At decision epoch t_0 the agent chooses an action/policy based on information \mathcal{I}_{t_0} and a predictive object Q (e.g., distribution, score, or control). Outcomes Y are realized later; verification arrives after a lag $\tau \geq 0$. The *action window* has length $\Delta > 0$.

Definition 1 (Ideal benchmark). *L_{ideal} is the minimal expected loss in a counterfactual world with zero verification lag ($\tau = 0$), unbounded window ($\Delta \rightarrow \infty$), and order-commuting operators (no reflexivity): the Bayes risk given true conditionals and full information.*

Operational note. In strongly reflexive systems where operator non-commutativity is fundamental, the counterfactual limit defining L_{ideal} may not have a unique physical realization. In such cases L_{ideal} serves as an idealized reference point; our decomposition (g_1, g_2, g_{12}) quantifies structural penalties relative to this reference, whether or not it corresponds to an attainable state in a given deployment.

Loss as regret. We assume a strictly proper scoring rule (or a convex surrogate) inducing a Bregman divergence D_Φ (Sec. 4). The *excess loss* (regret) w.r.t. the ideal is

$$L - L_{\text{ideal}} = \mathbb{E}[D_\Phi(P^* \| Q)],$$

where P^* denotes the ideal predictive object (true conditional law or Bayes-optimal score).

3 Constraints as Convex Sets

We encode two constraint families as closed convex sets in the prediction space.

Definition 2 (Latency-feasible set \mathcal{C}_λ). *For a latency parameter $\lambda \equiv \lambda(\tau, \Delta)$ (increasing in τ/Δ), the set \mathcal{C}_λ contains predictive objects Q that are measurable with respect to the restricted information σ -algebra available before decisions must be finalized.*

Definition 3 (Order-feasible set \mathcal{C}_ε). *For an order-sensitivity index $\varepsilon_* \geq 0$, the set \mathcal{C}_ε contains predictive objects Q achievable under operator orderings that satisfy specified reflexivity/disclosure constraints.*

The feasible region is $\mathcal{C}(\lambda, \varepsilon_*) = \mathcal{C}_\lambda \cap \mathcal{C}_\varepsilon$ (assumed nonempty).

Standing assumptions.

Assumption 1. \mathcal{C}_λ and \mathcal{C}_ε are nonempty, closed, and convex with respect to the affine structure of the prediction space; Φ is Legendre (strictly convex, essentially smooth).

Domain guidance. The convexity assumption is natural in:

- Probability forecast spaces (distributions form a convex set).
- Risk-neutral pricing (loss is linear in probabilities).
- Randomized policies (mixtures across episodes).

It may fail in:

- Hard sequential decisions with discrete action trees.
- Strongly path-dependent systems with irreversible state changes.

When convexity fails, we recommend using Section 8.2 (convexification) or Section A.3 (mirror geometry) with explicit tracking of δ_{relax} as part of Δ_{ncx} .

4 Bregman Geometry Preliminaries

Let Φ be a strictly convex, differentiable potential on a finite-dimensional open convex set \mathcal{X} . The Bregman divergence

$$D_\Phi(p\|q) = \Phi(p) - \Phi(q) - \langle \nabla\Phi(q), p - q \rangle$$

is nonnegative and equals zero iff $p = q$. For a nonempty, closed, convex set $\mathcal{C} \subseteq \mathcal{X}$ and any $p \in \mathcal{X}$, the *Bregman projection* is

$$\Pi_{\mathcal{C}}(p) \in \arg \min_{q \in \mathcal{C}} D_\Phi(p\|q).$$

Lemma 1 (Generalized Pythagorean inequality). *If $q^* = \Pi_{\mathcal{C}}(p)$, then for any $q \in \mathcal{C}$,*

$$D_\Phi(p\|q) \geq D_\Phi(p\|q^*) + D_\Phi(q^*\|q).$$

Proof. First-order optimality for q^* yields $\langle \nabla\Phi(q^*) - \nabla\Phi(q), q - q^* \rangle \geq 0$ for all $q \in \mathcal{C}$. Expanding $D_\Phi(p\|q) - D_\Phi(p\|q^*) - D_\Phi(q^*\|q)$ and applying this variational inequality gives the claim. \square

5 Main Result: Yis Structured Lower-Bound Decomposition

Define the *ideal* predictive object $p^* \equiv P^*$. Let

$$q_\varepsilon = \Pi_{\mathcal{C}_\varepsilon}(p^*), \quad q_{\lambda,\varepsilon} = \Pi_{\mathcal{C}_\lambda}(q_\varepsilon), \quad q_{\text{feas}} = \Pi_{\mathcal{C}(\lambda,\varepsilon_*)}(p^*).$$

Theorem 1 (Yis Structured Lower-Bound Decomposition). *Under Assumption 1, for any feasible $q \in \mathcal{C}(\lambda, \varepsilon_*)$,*

$$D_\Phi(p^*\|q) \geq \underbrace{D_\Phi(p^*\|q_\varepsilon)}_{g_2(\varepsilon_*)} + \underbrace{D_\Phi(q_\varepsilon\|q_{\lambda,\varepsilon})}_{g_1(\lambda)} + \underbrace{D_\Phi(q_{\lambda,\varepsilon}\|q)}_{g_{12}(\lambda, \varepsilon_*)}. \quad (1)$$

Consequently,

$$L - L_{\text{ideal}} \geq g_1(\lambda) + g_2(\varepsilon_*) + g_{12}(\lambda, \varepsilon_*),$$

where g_1 and g_2 are nonnegative and monotone in their respective arguments, and $g_{12} \geq 0$ by definition, with $g_{12} = 0$ if and only if $q_{\lambda,\varepsilon} = q_{\text{feas}}$ (orthogonality / commuting projections).

Proof. Apply Lemma 1 with $\mathcal{C} = \mathcal{C}_\varepsilon$ to get $D_\Phi(p^*\|q) \geq D_\Phi(p^*\|q_\varepsilon) + D_\Phi(q_\varepsilon\|q)$ for any $q \in \mathcal{C}_\varepsilon$ (in particular any feasible q). Then apply Lemma 1 again on $\mathcal{C} = \mathcal{C}_\lambda$ with $p = q_\varepsilon$ and $q \in \mathcal{C}_\lambda$ (again any feasible q qualifies), to obtain $D_\Phi(q_\varepsilon\|q) \geq D_\Phi(q_\varepsilon\|q_{\lambda,\varepsilon}) + D_\Phi(q_{\lambda,\varepsilon}\|q)$. Summing the two inequalities yields Eq. (1). Taking expectations over contexts (if any) transfers to $L - L_{\text{ideal}}$. Monotonicity:

if $\lambda_1 \leq \lambda_2$ then $\mathcal{C}_{\lambda_1} \supseteq \mathcal{C}_{\lambda_2}$, so the projection distance cannot decrease when shrinking the set; similarly for ε_\star . Finally $g_{12} \geq 0$ is $D_\Phi(q_{\lambda,\varepsilon} \| q) \geq 0$, with equality iff $q = q_{\lambda,\varepsilon}$. When the intersection projection equals the sequential projection ($q_{\text{feas}} = q_{\lambda,\varepsilon}$), the minimal feasible q attains $g_{12} = 0$. \square

Remark 1 (Interpretation). g_2 is the order penalty: how far the ideal is from any order-feasible object. g_1 is the latency penalty given we have already respected order constraints. g_{12} captures the residual gap between sequentially enforcing the constraints and jointly enforcing both—a nonnegative “interaction” that vanishes when projections commute.

Remark 2 (Selection and the decomposition). While g_{12} is selection-free by construction (see the definition on p. 1), g_1 and g_2 may depend on which projection $a \in \arg \min_{r \in \mathcal{C}_\varepsilon} D_\Phi(p^* \| r)$ is chosen when uniqueness fails. However, for any optimal feasible point $q^* \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda$, the total lower bound

$$g_1(\lambda) + g_2(\varepsilon_\star) + g_{12}(\lambda, \varepsilon_\star)$$

is selection-invariant, because it equals $D_\Phi(p^* \| q^*)$; both sides are defined as minima over the same feasible set.

6 Units, Calibration, and Monotonicity

The divergence D_Φ inherits the *loss units* of the underlying strictly proper score. Thus g_1, g_2, g_{12} are already in actionable units (regret, dollars, risk points, service-level debt).

Proposition 1 (Boundary conditions). $g_1(0) = 0, g_2(0) = 0$; if $\tau = 0$ or $\Delta \rightarrow \infty$, then $\lambda = 0$ thus $g_1 = 0$; if operators commute (no reflexivity/disclosure sensitivity), then $\varepsilon_\star = 0$ thus $g_2 = 0$.

Proposition 2 (Monotonicity). If λ increases (shrinking information or shortening windows), $g_1(\lambda)$ is nondecreasing; if ε_\star increases (stronger reflexivity/order sensitivity), $g_2(\varepsilon_\star)$ is nondecreasing.

7 When Does the Interaction Vanish?

Theorem 2 (Orthogonality / commuting projections). Suppose the constraint sets are Bregman-orthogonal at $q_{\lambda,\varepsilon}$, i.e., the normal cones $N_{\mathcal{C}_\lambda}(q_{\lambda,\varepsilon})$ and $N_{\mathcal{C}_\varepsilon}(q_{\lambda,\varepsilon})$ span complementary subspaces under the dual geometry. Then $q_{\lambda,\varepsilon} = \Pi_{\mathcal{C}(\lambda,\varepsilon)}(p^*)$ and $g_{12} = 0$.

Proof sketch. Under complementary normal cones, sequential KKT conditions at $q_{\lambda,\varepsilon}$ satisfy the joint projection KKT system for the intersection. Hence the sequential projection equals the joint projection, and $D_\Phi(q_{\lambda,\varepsilon} \| q_{\text{feas}}) = 0$. \square

Remark 3. Operationally: if enforcing order constraints does not perturb the gradient of the latency constraint at the solution (and vice versa), interaction disappears. Otherwise, the constraints “push” in coupled directions, inducing $g_{12} > 0$.

8 Extensions Beyond Convexity and Legendre Assumptions

8.1 Beyond Convexity and Legendre Assumptions

Scope. Yi’s Framework Yi’s Framework decomposes

$$L \geq L_{\text{ideal}} + g_1(\lambda) + g_2(\varepsilon_\star) + g_{12}(\lambda, \varepsilon_\star),$$

under closed convex feasibility and a Legendre potential Φ that induces the Bregman divergence $D_\Phi(\cdot \parallel \cdot)$. Real systems are often nonconvex and non-Legendre. We provide four robust paths that keep a valid lower bound and clear operational guidance.

Unified penalized form. In general deployments we will track

$$L \geq L_{\text{ideal}} + g_1(\lambda) + g_2(\varepsilon_\star) + g_{12}(\lambda, \varepsilon_\star) - \Delta_{\text{ncx}},$$

where $\Delta_{\text{ncx}} \geq 0$ aggregates nonconvexity/approximation penalties (e.g., convex-hull relaxation gaps, local-curvature remainder ζ , and empirical truncation), with $\Delta_{\text{ncx}} = 0$ under convex or Φ -geodesic-convex conditions.

Nonconvexity penalty decomposition. We write $\Delta_{\text{ncx}} := \delta_{\text{relax}} + \zeta + \delta_{\text{emp}} \geq 0$, where δ_{relax} is the convex-hull relaxation gap from replacing $(\mathcal{C}_\varepsilon, \mathcal{C}_\lambda)$ by $(\widehat{\mathcal{C}}_\varepsilon, \widehat{\mathcal{C}}_\lambda)$, ζ is the local curvature remainder from prox-regular/geodesic analysis, and δ_{emp} is the empirical truncation/monitoring component induced by the 2×2 or DR evaluation. Each term is either estimable or admits an operational upper bound reported with the lower bound.

8.2 Convex Relaxation, Always-Valid Conservative Bounds

Let $\widehat{\mathcal{C}}_\varepsilon := \text{cl conv}(\mathcal{C}_\varepsilon)$ and $\widehat{\mathcal{C}}_\lambda := \text{cl conv}(\mathcal{C}_\lambda)$ and define $\hat{a} = \Pi_{\widehat{\mathcal{C}}_\varepsilon}^\Phi(p^*)$, $\hat{b} = \Pi_{\widehat{\mathcal{C}}_\lambda}^\Phi(\hat{a})$, with

$$\widehat{g}_2 = D_\Phi(p^* \parallel \hat{a}), \quad \widehat{g}_1 = D_\Phi(\hat{a} \parallel \hat{b}), \quad \widehat{g}_{12} = \inf_{q \in \widehat{\mathcal{C}}_\varepsilon \cap \widehat{\mathcal{C}}_\lambda} D_\Phi(\hat{b} \parallel q).$$

Theorem (short). For any $q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda$, $L - L_{\text{ideal}} \geq \widehat{g}_2 + \widehat{g}_1 + \widehat{g}_{12}$. *Use:* safety & audits; transfers directly to OEC if L aligns with D_Φ (proper scoring/calibrated proxy).

Remark 4 (Conservativeness of convexification). *Since $\widehat{\mathcal{C}}_\varepsilon \supseteq \mathcal{C}_\varepsilon$ and $\widehat{\mathcal{C}}_\lambda \supseteq \mathcal{C}_\lambda$, we also have $\widehat{\mathcal{C}}_\varepsilon \cap \widehat{\mathcal{C}}_\lambda \supseteq \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda$. Each of $\widehat{g}_2, \widehat{g}_1, \widehat{g}_{12}$ is defined via a minimization over these convexified sets, so the convexified geometric gap $\widehat{g}_2 + \widehat{g}_1 + \widehat{g}_{12}$ cannot exceed the original gap $g_2 + g_1 + g_{12}$ when the latter is well-defined. Thus the convexified lower bound is smaller (more conservative) while remaining valid for all $q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda$.*

8.3 Local Analysis under Prox-Regularity

With Φ α -strongly convex and smooth near $\{p^*, a, b\}$, and $\mathcal{C}_\varepsilon, \mathcal{C}_\lambda$ prox-regular locally,

$$L - L_{\text{ideal}} \geq g_2 + g_1 + g_{12}^{\text{rob}} - \zeta, \quad \zeta \geq 0.$$

A practical upper bound:

$$\boxed{\zeta \leq \frac{\rho}{\alpha} (D_\Phi(p^* \parallel a) + D_\Phi(a \parallel b)) + c \cdot \widehat{\kappa} (\delta \lambda \delta \varepsilon)^2},$$

where ρ is weak nonconvexity (prox-regular modulus), and $\widehat{\kappa}$ a local curvature density measured by a micro 2×2 toggle. *Unified view:* the remainder ζ is subsumed into Δ_{ncx} .

8.4 Assumption-Light Empirical Decomposition

Define four regimes $(L_{00}, L_{01}, L_{10}, L_{11})$ via toggled/staggered latency & order. Estimators:

$$\hat{g}_2 = L_{01} - L_{00}, \quad \hat{g}_1 = L_{11} - L_{01}, \quad \hat{g}_{12} = L_{11} - L_{01} - L_{10} + L_{00}, \quad [\hat{g}_{12}]_+ = \max\{0, \hat{g}_{12}\}.$$

Lemma (short). $L - L_{\text{ideal}} \geq \hat{g}_2 + \hat{g}_1 + [\hat{g}_{12}]_+$. *Practice:* estimate L_{ij} by DR/IPW (selection) + IPCW (censoring); report ESS & clipping%; clustered CIs.

Weights, truncation, and ESS. Inverse-propensity and censoring weights can be heavy-tailed; we therefore enforce explicit truncation and reporting rules. Fix a truncation threshold $c \geq 1$ and let w denote the product of selection and censoring weights. A standard bound for the truncation bias is

$$|\text{Bias}| \lesssim \mathbb{P}(w > c) \cdot \sup |Y| + \mathcal{O}(c^{-1}),$$

so we always report: (i) the chosen threshold c , (ii) the fraction of mass with $w > c$ (clipping%), (iii) the effective sample size $\text{ESS} = (\sum_i w_i)^2 / \sum_i w_i^2$, and (iv) the untruncated estimate side by side with the truncated one. In deployments where ESS falls below a minimum threshold (e.g. 100–200), we recommend using the estimates only for monitoring rather than for hard guarantees. We also adopt doubly robust estimation with cross-fitting, stabilized weights, and optionally TMLE/DR-learner style targeted regressions to improve numerical stability; the residual uncertainty is absorbed into the empirical component of Δ_{ncx} .

Unified view: the truncation control contributes to Δ_{ncx} (empirical truncation component).

8.5 Φ -Geodesic Convexity (mirror/g-convex) Path

Some sets are not Euclidean-convex but are convex in the mirror geometry induced by a Legendre Φ . For closed g-convex $\mathcal{C}_\varepsilon, \mathcal{C}_\lambda$, mirror projections are unique and satisfy a mirror Pythagorean identity, yielding

$$L - L_{\text{ideal}} \geq g_2 + g_1 + g_{12}^\Phi \quad (\text{no } \zeta).$$

Use: choose Φ aligned with the domain (entropy on simplex; quadratic on subspaces; log-partition for exponential families). *Unified view:* under Φ -geodesic convexity, $\Delta_{\text{ncx}} = 0$ and the clean additive law is recovered.

8.6 Guidance Matrix (When to Use Which Path)

Scenario	Path	Output	KPI
Safety & audit	Section 8.2 + Section 8.4	Conservative bound + empirical check	$\hat{\kappa} \downarrow, \text{ESS} \uparrow$
Performance tuning	Section 8.3 + Section 8.4	Tight bound, ζ monitor, A/B confirm	$\zeta \downarrow, \text{OEC} \uparrow$
Routine monitor	Section 8.4	Real-time components	$[\hat{g}_{12}]_+, \text{clipping}\%$
Geometry-aligned	Section 8.5 + Section 8.4	Clean additive (mirror)	$[\hat{g}_{12}]_+ \downarrow$
Diagnosis	Section 8.3 (ζ)	Geometry vs. interaction bottlenecks	$\hat{\kappa} \downarrow, \zeta \downarrow$

Remark 5 (Semantics of penalized lower bound). *The nonconvexity penalty only loosens the lower bound. A safe report is $\text{LB}_{\text{safe}} = \max\{0, g_1 + g_2 + g_{12} - \Delta_{\text{ncx}}\}$. When $\Delta_{\text{ncx}} > g_1 + g_2 + g_{12}$, the bound is vacuous but still valid; then default to Section 8.4/Section 8.5 to reduce the penalty.*

9 Relation to Classical Decompositions

Biasvariance (static). For point estimation with quadratic loss and a static target, bias and variance are orthogonal in expectation, yielding no cross-term. That corresponds to $\varepsilon_* = 0$ (no reflexivity) and $\lambda = 0$ (no lag), hence $g_1 = g_2 = g_{12} = 0$ relative to the static Bayes risk. Yi's Framework generalizes to dynamic, reflexive, finite-window pipelines where the additive penalties are strictly positive.

CAP-like limits (latency extreme). Under effective partition/communication failure, $\tau \rightarrow \infty \Rightarrow \lambda \rightarrow \infty$ and g_1 dominates; the theory subsumes distributed and experimental constraints as special cases of the latency term.

10 Estimation of g_1, g_2, g_{12}

Let $L(\cdot)$ denote measured loss/regret under controlled regimes.

Two-stage projection emulation. Construct four regimes: Unconstrained, Order-only, Latency-only, Both. Let $L_{00}, L_{01}, L_{10}, L_{11}$ be corresponding losses above L_{ideal} (estimated via gold references or extrapolation). Define

$$\hat{g}_2 = L_{01}, \quad \hat{g}_1 = L_{11} - L_{01}, \quad \hat{g}_{12} = L_{11} - L_{01} - L_{10} + L_{00} \geq 0.$$

The nonnegativity of \hat{g}_{12} is implied by Eq. (1) when the regimes emulate sequential and joint projections. Confidence intervals follow from block bootstraps or delta methods.

Interference, reflexivity, and SUTVA. The 2×2 layout is an *emulation* of the two-stage projections, not a literal switch that turns reflexivity or feedback off. In practice we approximate the ideal regimes by: (i) using shadow evaluation or off-policy logging windows in which the deployed policy is held fixed while outcomes are recorded, so that current decisions do not feed back into the state during the evaluation window; (ii) randomizing at the level of clusters, buckets, or episodes so that interference is allowed within clusters but assumed negligible across clusters (partial interference); and (iii) using encouragement designs where an assignment Z nudges the use of a constrained or unconstrained policy while sequential order acts as a mediator. These design choices make the working SUTVA/partial-interference assumptions explicit.

SUTVA sensitivity. We introduce $\Delta_{\text{SUTVA}} \geq 0$ as a sensitivity radius that upper-bounds possible bias in \hat{g}_{12} due to violations of these assumptions. Alongside the point estimate we report an interval

$$\hat{g}_{12} \in [\hat{g}_{12}^{\text{naive}} - \Delta_{\text{SUTVA}}, \hat{g}_{12}^{\text{naive}} + \Delta_{\text{SUTVA}}],$$

with Δ_{SUTVA} calibrated by domain knowledge (cluster sizes, leakage rates) or simulation. Correspondingly, the empirical lower bound may be further relaxed as

$$L - L_{\text{ideal}} \geq \hat{g}_2 + \hat{g}_1 + [\hat{g}_{12}]_+ - \Delta_{\text{ncx}} - \Delta_{\text{SUTVA}},$$

when a conservative accounting for interference is required; in this view, Δ_{SUTVA} can be treated as part of the empirical component of Δ_{ncx} .

Continuous calibration. If λ and ε_* are graded, fit shape-constrained monotone regressions $g_1(\lambda)$, $g_2(\varepsilon_*)$ (e.g., isotonic or convex regression) and report g_{12} as the residual consistent with subadditivity bounds.

11 Worked Example: Gaussian Control with Lag and Reflexivity

Consider $Y \in \mathbb{R}$, $Y \sim \mathcal{N}(\mu, \sigma^2)$. The agent outputs $Q = \mathcal{N}(m, v)$, scored by log-loss; then $D_\Phi(P^* \| Q) = \frac{(m-\mu)^2}{2v} + \frac{1}{2} \left(\frac{\sigma^2}{v} - 1 - \log \frac{\sigma^2}{v} \right)$.

Latency constraint. Suppose m must be formed from a *lagged proxy* $\tilde{\mu}$ with $\tilde{\mu} \sim \mathcal{N}(\mu, \sigma_\lambda^2)$ independent of μ and $v = \sigma^2$. Then $g_1(\lambda) = \mathbb{E}\left[\frac{(\tilde{\mu}-\mu)^2}{2\sigma^2}\right] = \frac{\sigma_\lambda^2}{2\sigma^2}$, monotone in the proxy MSE induced by lag.

Order constraint (reflexivity). Suppose actions shift μ by δa and exposure mixes distributions (disclosure), limiting feasible (m, v) to a convex set that shrinks with ε_* . The projection onto this set increases the mean error by $\Delta m(\varepsilon_*)$, yielding $g_2(\varepsilon_*) = \frac{(\Delta m(\varepsilon_*))^2}{2\sigma^2}$.

Interaction. If the proxy noise correlates with action-induced shifts (e.g., the same channels both delay labels and cause exposure), the joint-feasible projection is stricter than the sequential one, generating $g_{12} > 0$. This toy model makes the abstract terms computable and illustrates monotonicity.

12 Proof Details

We collect key proofs for completeness.

Lemma 2 (Existence/uniqueness of projections). *Under Assumption 1, $\Pi_C(p)$ exists and is unique for D_Φ with Legendre Φ .*

Proof. $D_\Phi(p\|\cdot)$ is strictly convex and lower semicontinuous on closed convex C ; a unique minimizer exists by standard convex analysis. \square

Lemma 3 (Two-stage projection bound). *Let \mathcal{A}, \mathcal{B} be closed convex sets with nonempty intersection, $p \in \mathcal{X}$, and define $a = \Pi_{\mathcal{A}}(p)$, $b = \Pi_{\mathcal{B}}(a)$, $c = \Pi_{\mathcal{A} \cap \mathcal{B}}(p)$. Then for any $q \in \mathcal{A} \cap \mathcal{B}$,*

$$D_\Phi(p\|q) \geq D_\Phi(p\|a) + D_\Phi(a\|b) + D_\Phi(b\|q).$$

Proof. Apply Lemma 1 with (p, \mathcal{A}) to split off $D_\Phi(p\|a)$ and with (a, \mathcal{B}) to split off $D_\Phi(a\|b)$; the remainder is $D_\Phi(b\|q) \geq 0$. \square

Proof of Theorem 1. Identify $\mathcal{A} = \mathcal{C}_\varepsilon$, $\mathcal{B} = \mathcal{C}_\lambda$, $p = p^*$, $a = q_\varepsilon$, $b = q_{\lambda, \varepsilon}$, q feasible; then invoke the lemma and take expectations over contexts. \square

Proof of Theorem 2. Under complementary normal cones, sequential KKT multipliers jointly satisfy the intersection KKT system; thus $b = c$, implying $g_{12} = D_\Phi(b\|c) = 0$. \square

13 Operational Playbook (Synopsis)

Measure & report. Instrument τ (lag) and Δ (window) to derive λ . Measure order-sensitivity ε_* via counterfactual replay and order perturbations. Publish $(L_{\text{ideal}}, \hat{g}_1, \hat{g}_2, \hat{g}_{12})$ with CIs.

Route by smallest structural penalty. When λ dominates, prefer many *reversible* steps (fast VA). When ε_* dominates, prefer conservative, well-verified PV. Use PA when models are robust and actions reversible. Always maintain rollbacks and blast-radius limits.

Invest to shrink penalties. Parallel/stream verification lowers g_1 ; partial disclosure, decoupled interfaces, and shadow evaluation lower g_2 . Seek orthogonality to reduce g_{12} .

14 Discussion and Limitations

Yi's Framework provides a geometry that unifies disparate "impossibility" phenomena under time and order constraints. The main limitation is *calibration*: mapping λ, ε_* to concrete regimes requires domain-specific design. The framework diagnoses structural floors; it complements, not replaces, ethical, legal, and human-centric constraints.

15 Conclusion

We proved a single inequality that decomposes unavoidable loss into time, order, and interaction components via Bregman projections. This converts abstract trade-offs into measurable, optimizable quantities, enabling principled routing and investment across domains. In deployments with nonconvex feasibility or approximate projections, we adopt the penalized additive bound with $\Delta_{\text{ncx}} \geq 0$ and report LB_{safe}; under convex or Φ -geodesic-convex regimes, $\Delta_{\text{ncx}} = 0$ and the clean additive law holds.

Keywords. Yi's Framework; Bregman divergence; projection; latency; noncommutativity; reflexivity; lower bound; diagnostics.

A Appendix A: Robust Extensions

A.1 Convex Relaxation, Always-Valid Conservative Bounds

Construction. Let $\widehat{\mathcal{C}}_\varepsilon := \text{cl conv}(\mathcal{C}_\varepsilon)$ and $\widehat{\mathcal{C}}_\lambda := \text{cl conv}(\mathcal{C}_\lambda)$. Define relaxed projections

$$\hat{a} := \Pi_{\widehat{\mathcal{C}}_\varepsilon}^\Phi(p^*), \quad \hat{b} := \Pi_{\widehat{\mathcal{C}}_\lambda}^\Phi(\hat{a}),$$

and relaxed penalties

$$\hat{g}_2 := D_\Phi(p^* \| \hat{a}), \quad \hat{g}_1 := D_\Phi(\hat{a} \| \hat{b}), \quad \hat{g}_{12} := \inf_{q \in \widehat{\mathcal{C}}_\varepsilon \cap \widehat{\mathcal{C}}_\lambda} D_\Phi(\hat{b} \| q).$$

Theorem 3 (Convexified lower bound — always valid). *For any realized $q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda$ (hence $q \in \widehat{\mathcal{C}}_\varepsilon \cap \widehat{\mathcal{C}}_\lambda$) and any Legendre Φ ,*

$$D_\Phi(p^* \| q) \geq \hat{g}_2 + \hat{g}_1 + \hat{g}_{12}.$$

Consequently, $L - L_{\text{ideal}} \geq \hat{g}_2 + \hat{g}_1 + \hat{g}_{12}$.

Proof sketch. $\widehat{\mathcal{C}}_\varepsilon, \widehat{\mathcal{C}}_\lambda$ are closed convex; the Bregman Pythagorean identity gives $D_\Phi(p^* \| q) \geq D_\Phi(p^* \| \hat{a}) + D_\Phi(\hat{a} \| \hat{b}) + D_\Phi(\hat{b} \| q)$ for all $q \in \widehat{\mathcal{C}}_\varepsilon \cap \widehat{\mathcal{C}}_\lambda$. Taking \inf_q on the RHS yields the claim. Any actual feasible q lies in the relaxed intersection, so the bound is conservative. \square

Remark 6 (OEC alignment). *When L is aligned with D_Φ (proper scoring or calibrated proxy), Theorem 3 transfers directly to the operational objective; otherwise read it as a geometric audit bound.*

A.2 Local Prox-Regular Analysis with Curvature Penalty

Assumptions. (i) Φ is α -strongly convex and L -smooth near $\{p^*, a, b\}$. (ii) $\mathcal{C}_\varepsilon, \mathcal{C}_\lambda$ are prox-regular near a, b (local single-valued projections; bounded set curvature).

Proposition 3 (Local robust decomposition with curvature penalty). *Let $a \in \arg \min_{u \in \mathcal{C}_\varepsilon} D_\Phi(p^* \| u)$, $b \in \arg \min_{v \in \mathcal{C}_\lambda} D_\Phi(a \| v)$ and*

$$g_{12}^{\text{rob}} := \inf_{q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda} D_\Phi(b \| q).$$

Then there exists $\zeta = \zeta(\alpha, L, \text{curvature at } a, b) \geq 0$ such that for all feasible $q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda$,

$$D_\Phi(p^* \| q) \geq \underbrace{D_\Phi(p^* \| a)}_{g_2} + \underbrace{D_\Phi(a \| b)}_{g_1} + g_{12}^{\text{rob}} - \zeta.$$

Hence $L - L_{\text{ideal}} \geq g_2 + g_1 + g_{12}^{\text{rob}} - \zeta$.

Operable upper bound for ζ . With weak nonconvexity ρ (prox-regular/weakly-convex modulus) and Φ -strong convexity α ,

$$\zeta \leq \frac{\rho}{\alpha} (D_\Phi(p^* \| a) + D_\Phi(a \| b)) + c \cdot \hat{\kappa} (\delta \lambda \delta \varepsilon)^2,$$

where c is a problem-dependent constant and $\hat{\kappa}$ is a local curvature density from a 2x2 micro-experiment.

Remark 7 (Semantics of penalized bound). *The penalty only loosens the lower bound. A safe report is $\text{LB}_{\text{safe}} = \max\{0, g_1 + g_2 + g_{12} - \zeta\}$. When ζ exceeds $g_1 + g_2 + g_{12}$, the bound is vacuous but valid; switch to Section 8.4/Section 8.5 to reduce the penalty.*

A.3 Assumption-Light Empirical 2x2 Decomposition

Four regimes (toggle/staggered design).

	order/NC on	order/NC off
latency on	L_{11}	L_{01}
latency off	L_{10}	L_{00}

Estimators. $\hat{g}_2 = L_{01} - L_{00}$, $\hat{g}_1 = L_{11} - L_{01}$, $\hat{g}_{12} = L_{11} - L_{01} - L_{10} + L_{00}$.

Lemma 4 (Robust lower bound). *With $[\hat{g}_{12}]_+ := \max\{0, \hat{g}_{12}\}$, $L - L_{\text{ideal}} \geq \hat{g}_2 + \hat{g}_1 + [\hat{g}_{12}]_+$.*

Practice. Estimate L_{ij} using DR/IPW (selection) and IPCW (censoring); report effective sample size (ESS) and weight-clipping rates (e.g., 99–99.5%). Provide clustered/bootstrap CIs. Use current $(\hat{g}_1, \hat{g}_2, [\hat{g}_{12}]_+)$ to route, and invest to reduce them over horizons.

A.4 Φ -Geodesic Convexity (Mirror/g-Convex Path)

Mirror geometry. Let $\nabla\Phi$ be the mirror map and $\nabla\Phi^*$ its inverse. The Φ -geodesic segment between x, y is

$$\Gamma_\Phi(x, y; t) := \nabla\Phi((1-t)\nabla\Phi(x) + t\nabla\Phi(y)), \quad t \in [0, 1].$$

Definition 4 (Φ -geodesic convexity). A set $C \subset X$ is Φ -geodesically convex (g-convex) if $\Gamma_\Phi(x, y; t) \in C$ for all $x, y \in C, t \in [0, 1]$.

Projection and Pythagorean in mirror geometry. For closed g-convex C , define $\Pi_C^\Phi(x) := \arg \min_{y \in C} D_\Phi(x \| y)$. Under Legendre Φ , the mirror projection is unique and satisfies a mirror Pythagorean inequality:

$$D_\Phi(x \| z) \geq D_\Phi\left(x \middle\| \Pi_C^\Phi(x)\right) + D_\Phi\left(\Pi_C^\Phi(x) \middle\| z\right), \quad \forall z \in C.$$

Theorem 4 (Lower bound under Φ -geodesic convexity). Suppose \mathcal{C}_ε and \mathcal{C}_λ are closed and g-convex. With $a = \Pi_{\mathcal{C}_\varepsilon}^\Phi(p^*)$, $b = \Pi_{\mathcal{C}_\lambda}^\Phi(a)$ and $g_{12}^\Phi := \inf_{q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda} D_\Phi(b \| q)$, we have

$$D_\Phi(p^* \| q) \geq D_\Phi(p^* \| a) + D_\Phi(a \| b) + g_{12}^\Phi, \quad \forall q \in \mathcal{C}_\varepsilon \cap \mathcal{C}_\lambda,$$

hence $L - L_{\text{ideal}} \geq g_2 + g_1 + g_{12}^\Phi$ without the curvature penalty ζ .

Checks and examples. Pick Φ aligned with the domain: negative entropy on the simplex (probabilities), Mahalanobis/quadratic for Euclidean subspaces, log-partition families for exponential families. Empirically check approximate g-convexity by: (i) testing whether $\Gamma_\Phi(x, y; t)$ stays feasible for random x, y ; (ii) testing near-orthogonality of normal directions of $\mathcal{C}_\varepsilon, \mathcal{C}_\lambda$ under the $G = \nabla^2\Phi$ metric.

A.5 Reporting Convention and Diagnostics

We report: (i) $\text{LB}_{\text{safe}} = \max\{0, g_1 + g_2 + g_{12} - \Delta_{\text{ncx}}\}$; (ii) a penalty ratio

$$r := \frac{\Delta_{\text{ncx}}}{g_1 + g_2 + g_{12} + \varepsilon}, \quad \varepsilon > 0 \text{ small},$$

with heuristics: $r \leq 0.5$ informative; $r > 1$ vacuous (default to Section 8.4/Section 8.5). We also report ESS, clipping%, and CIs (clustered bootstrap).

A.6 Guidance Matrix (Full)

Scenario	Recommended Path	Output	KPI
Safety & compliance audit	A.1 (convexify) + A.3	Conservative bound + empirical verification	$\hat{\kappa} \downarrow, \text{ESS} \uparrow$
Performance optimization	A.2 (local) + A.3	Tight bound, ζ monitoring, A/B confirmation	$\zeta \downarrow, \text{OEC} \uparrow$
Routine monitoring	A.3 only	Real-time dashboard components	$[\hat{g}_{12}]_+, \text{clipping}\%$
Geometry-aligned domains	A.4 (g-convex) + A.3	Clean additive bound in mirror geometry	$[\hat{g}_{12}]_+ \downarrow$
Root-cause diagnosis	A.2 ζ analysis	Identify geometric vs. interaction bottlenecks	$\hat{\kappa} \downarrow, \zeta \downarrow$